

LLM Finetune with



unsloth

Why need this?

- Make response better
- Add more knowledge

Text
sample:

LLMs learn to predict one word at a time

LLMs learn to predict one word at a time

LLMs learn to predict one word at a time

LLMs learn to predict one word at a time

LLMs learn to predict one word at a time

LLMs learn to predict one word at a time

LLMs learn to predict one word at a time

LLMs learn to predict one word at a time

The LLM
can't access
words past
the target


Target to
predict

Input the
LLM
receives

Can we train LLM?

The Extreme Cost Of Training AI Models

Estimated cost of training selected AI models (in million U.S. dollars), by different calculation models



Rounded numbers. Excludes staff salaries that can make up 29-49% of final cost (including equity)


Source: Epoch AI



statista

OpenAI is estimated to lose \$5B in 2024

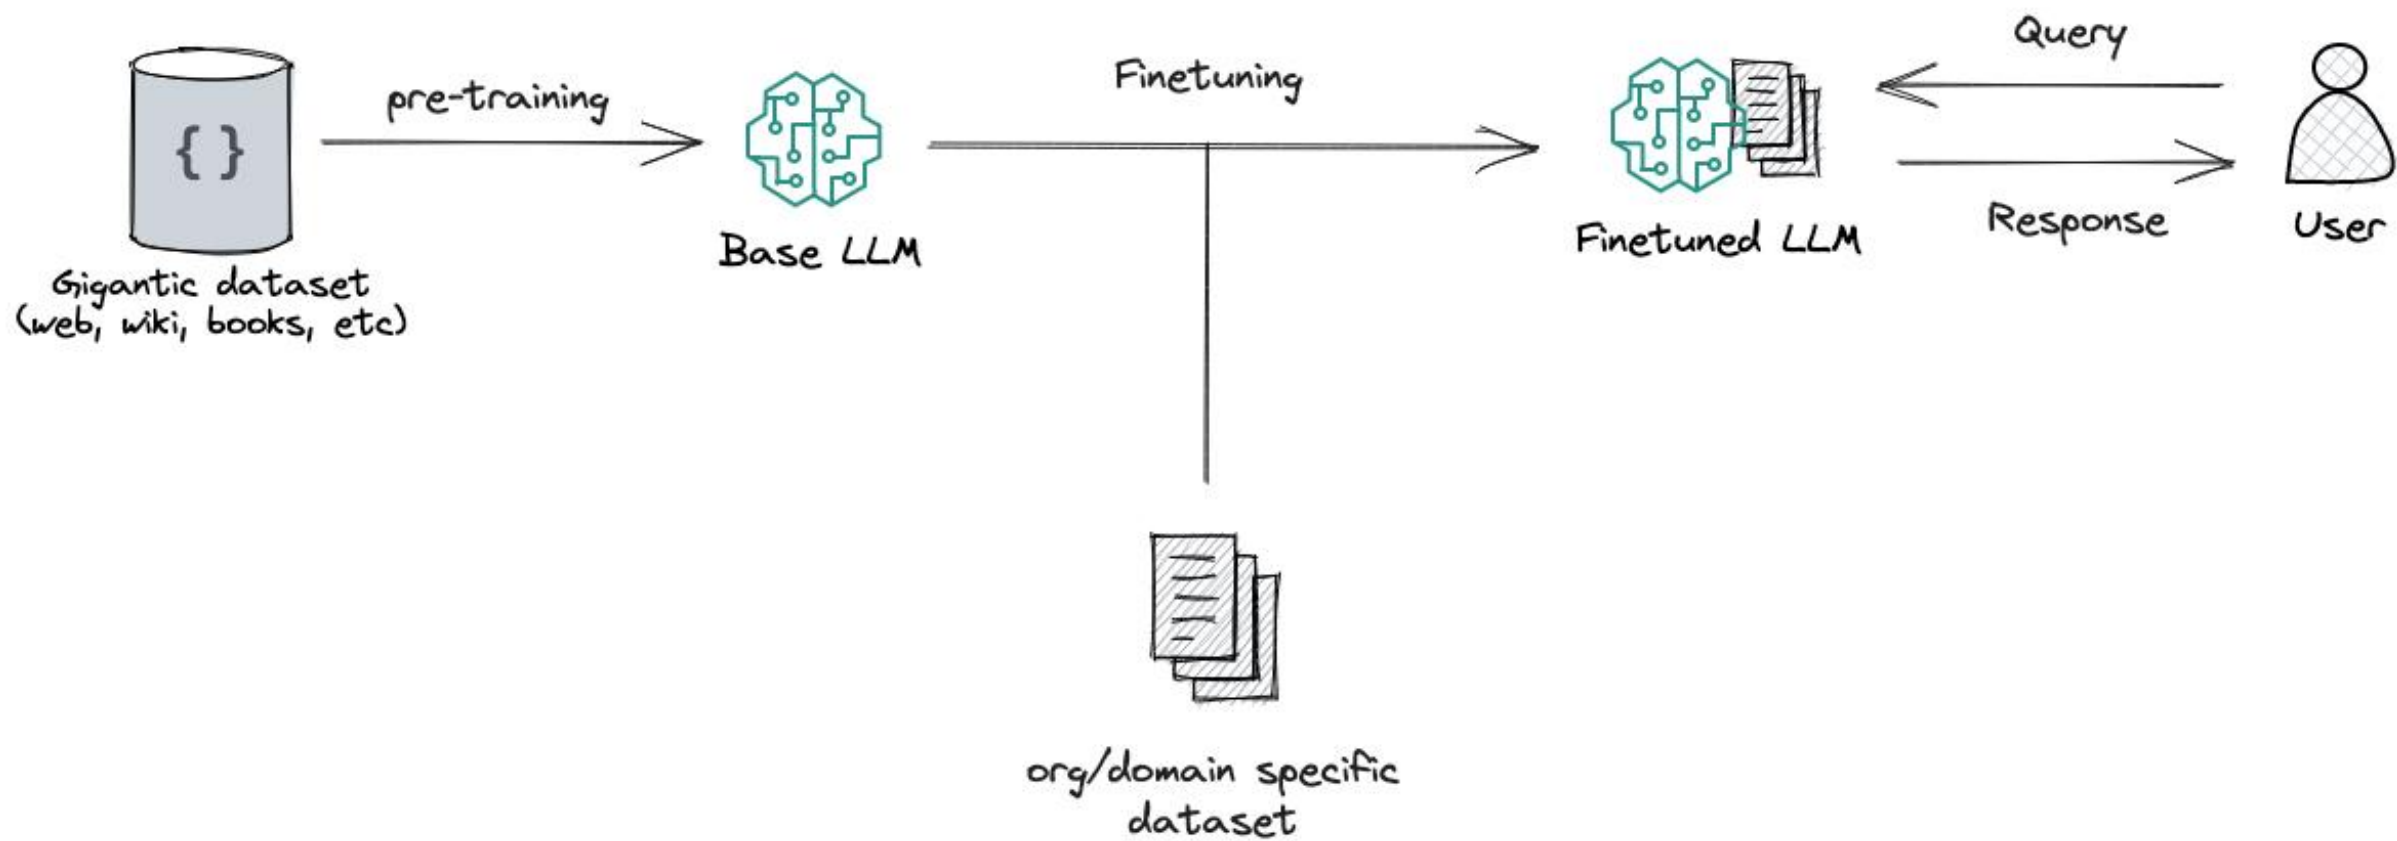
X Post



Source: The Information, internal analysis.

Understand Fine-tuning

- Fine-tuning an LLM customizes its behavior, enhances + injects knowledge, and optimizes performance for domains/specific tasks. For example:
 - GPT-4 serves as a base model; however, OpenAI fine-tuned it to better comprehend instructions and prompts, leading to the creation of ChatGPT-4 which everyone uses today.
 - DeepSeek-R1-Distill-Llama-8B is a fine-tuned version of Llama-3.1-8B. DeepSeek utilized data generated by DeepSeek-R1, to fine-tune Llama-3.1-8B. This process, known as distillation (a subcategory of fine-tuning), injects the data into the Llama model to learn reasoning capabilities.



Base Model

(davinci (GPT-3))

Prompt: Tell me how to fine-tune a model

Completion:

How can I control the complexity of a model?

How do I know when my model is done?

How do I test a model?

How do I define a task for my model?

How do I scale a model?

How do I make a probabilistic model?

How do I start using a model once it's complete?

How do I cope with missing data?

How do I make my modeling efforts sustainable?

Fine-tuned Model

(text-davinci-003)

Prompt: Tell me how to fine-tune a model

Completion:

Fine-tuning a model involves adjusting the parameters of a pre-trained model in order to make it better suited for a given task. There are generally three steps to fine-tuning a model:


1. Select a base model - Select a pre-trained deep learning model that has been trained on a large dataset.
2. Adjust parameters - Adjust parameters of the pre-trained model to better suit the desired task. This may include changing the number of layers, adjusting learning rate, adding regularization, or tweaking the optimizer.
3. Train the model - Train the new model on the desired dataset. The amount of data and the amount of training required will depend on the task and the model.

Lora & QLora

- LoRA: Fine-tunes small, trainable matrices in 16-bit without updating all model weights.
- QLoRA: Combines LoRA with 4-bit quantization to handle very large models with minimal resources.

Full Fine-Tuning


16-bit precision



- ✓ Best performance
- ✗ Very high VRAM usage

LoRA


16-bit precision



- ✓ Quick training
- ✗ Still costly

QLoRA

4-bit precision



- ✓ Low VRAM usage
- ✗ Degrades performance

What is unsloth?

Easily finetune & train LLMs
Get *faster* with unsloth

Unsloth is not just a library; it's a technological symphony orchestrated for the fine-tuning and training of large language models (LLMs).

Pros of Unsloth

Speed Redefined : Unsloth boasts a staggering 30x increase in training speed. Alpaca, a benchmark task, now takes merely 3 hours instead of the conventional 85. This acceleration is a testament to Unsloth's commitment to efficiency and productivity.

Memory Efficiency: A game-changer in the memory domain, Unsloth promises a 60% reduction in memory usage. This not only enables the handling of larger batches but also ensures a seamless fine-tuning process without compromising on performance.

Accuracy Amplified: The authors proudly declare a 0% loss in accuracy, with an additional option for a +20% increase in accuracy using their MAX offering. This commitment to maintaining and elevating accuracy levels sets Unsloth apart in the competitive landscape.

Hardware Compatibility: Unsloth extends its reach by supporting NVIDIA, Intel, and AMD GPUs. This inclusivity ensures accessibility to a wide array of hardware configurations, making it a versatile choice for developers across different platforms.

Step to Finetune

1. Choose the Right Model + Method
2. Prepare Your Dataset: You will need to create a dataset usually with 2 columns - question and answer. The quality and amount will largely reflect the end result of your fine-tune so it's imperative to get this part right.
3. Select training parameters
4. Train, save weight.
5. Test the finetuned model.


Step to Finetune

Understand Model Parameters:


- Leave almost alone
- Understand some of them:
 - Learning rate:
 - Higher Learning Rates: Faster training, reduces overfitting just make sure to not make it too high as it will overfit
 - Lower Learning Rates: More stable training, may require more epochs.
 - Typical Range: $1e-4$ (0.0001) to $5e-5$ (0.00005).
 - Epochs: 1-3 to avoid overfitting.

```
model = FastLanguageModel.get_peft_model(  
    model,  
    r = 16, # Choose any number > 0 ! Suggested 8, 16, 32, 64, 128  
    target_modules = ["q_proj", "k_proj", "v_proj", "o_proj",  
                     "gate_proj", "up_proj", "down_proj",],  
    lora_alpha = 16,  
    lora_dropout = 0, # Supports any, but = 0 is optimized  
    bias = "none",    # Supports any, but = "none" is optimized  
    # [NEW] "unsloth" uses 30% less VRAM, fits 2x larger batch sizes!  
    use_gradient_checkpointing = "unsloth", # True or "unsloth" for very long context  
    random_state = 3407,  
    use_rslora = False, # We support rank stabilized LoRA  
    loftq_config = None, # And LoftQ  
)
```


Overfitting vs Underfitting



Under-fitting
(too simple to
explain the variance)



Appropriate-fitting



Over-fitting
(forcefitting--too
good to be true) 

Train

- **per_device_train_batch_size = 2** – Increase for better GPU utilization but beware of slower training due to padding. Instead, increase gradient_accumulation_steps for smoother training.
- **gradient_accumulation_steps = 4** – Simulates a larger batch size without increasing memory usage.
- **max_steps = 60** – Speeds up training. For full runs, replace with num_train_epochs = 1 (1–3 epochs recommended to avoid overfitting).
- **learning_rate = 2e-4** – Lower for slower but more precise fine-tuning. Try values like 1e-4, 5e-5, or 2e-5.

Test

- In order to evaluate, you could do manually evaluation by just chatting with the model and see if it's to your liking.
- You can also enable evaluation for Unsloth, but keep in mind it can be time-consuming depending on the dataset size.
- To speed up evaluation you can: reduce the evaluation dataset size or set **evaluation_steps = 100**.

Handson

- Neet GPU \geq 24GB VRAM
- Linux: Ubuntu

